

Distributed Mutual Exclusion Algorithms: A Comparison of Central Server, Ring Token and Multicast

COMPSYS 725: Distributed Cyber-Physical Systems

Matt Eden

Department of Electrical, Computer and Software Engineering

University of Auckland

Auckland, New Zealand

mede607@aucklanduni.ac.nz

I. MUTUAL EXCLUSION ALGORITHMS

A. Overview

Mutual exclusion is generally concerned with preventing interference and ensuring consistency with resource access. For operating systems, this can be managed with relative ease. However, distributed systems do not have the shared variables or facilities that would be supplied by a single local kernel, so a different approach is required. This is where distributed mutual exclusion algorithms come into play, and a few of these are discussed in the following subsections.

There are a few key assumptions made in this report that should be acknowledged. The system being considered is comprised of N processes p_i with $i = 1, 2, 3 \dots N$. These process do not share variables, but access common resources. These resources are contained within a single critical section, and the system as a whole is asynchronous. The failure of processes is not considered, with message delivery being reliable so that any message sent is guaranteed to be delivered eventually and delivered exactly once.

B. Comparing Algorithms

The requirements for mutual exclusion considered in this report are:

- Safety [ME1]
- Liveness [ME2]
- Fairness [ME3]

Safety is concerned with the idea that at any instant in time only one process may occupy the critical section. *Liveness* is concerned with the idea that any and all requests to enter the critical section will eventually succeed. *Fairness* is concerned with how processes enter the critical section according to the order of requests. In this report I consider the ‘happened-before’ approach, wherein if one request to enter the critical section ‘happened-before’ another, then the entry is granted in that order. In addition to satisfying the above, performance of any one algorithm is evaluated against a set of fixed criteria, defined as follows.

- Consumed Bandwidth
- Client Delay
- Synchronisation Delay

The *Consumed Bandwidth* is proportional to the number of the messages sent in `entry()` and `exit()` (as defined previously). *Client Delay* is the time taken by a process at each `entry()` and `exit()`. *Synchronisation Delay* is simply the time difference between one process’ `exit()` and another’s `entry()`.

C. Central Server

The simplest algorithm described in this report consists of a central server managing requests for mutual exclusion from several processes. The overarching idea is that a token is held by this central server, loaned out to a process that has requested to enter the critical section and then retrieved once that process leaves the critical section. The central server maintains a record of requests it has received, and allocates tokens accordingly. The resolution of a request involves removing said request from the queue and allocating the token to that process. This workflow is illustrated in Figure 1, where p_2 is making a request to the server for the token, p_3 is releasing the token it held back to the server and p_4 is being allocated the token by the server as. Infer from the diagram that p_3 had previously made a request and obtained the token from the central server, hence why it is now releasing the token back to the server. Similarly, p_4 must have previously made a request to the central server as it is higher in the queue than p_2 .

This algorithm satisfies ME1 because, due to the inherent nature of the design, only one process can enter the critical section at a time. This is a direct result of the central server only allocating a single token, as when one process holds the token it cannot be allocated to any other process. ME2 is also satisfied by this algorithm, as a simple first-in first-out (FIFO) queue is used to order the requests by processes. Such an approach ensures the eventual allocation of a token to any process that requests it.

However, ME3 is not satisfied by the central server algorithm. The reason being that the ordering of requests is based purely on when the central server receives the request, not when the request itself is sent. As such, there is no guarantee that the ‘happened-before’ relationship is satisfied.

D. Ring Token

This approach does away with the dependence on a third-party element as seen in the central server algorithm, instead choosing to arrange processes in a logical ring. In doing so, the requirement is introduced that each process possesses a communication channel with the next. The token-based approach to allowing access to the critical section continues, with the token passed along the ring as illustrated in Figure 2. Upon receiving a token, a process may take one of two actions. Either it (a) enters the critical section or (b) passes the token to the next process. If a process has requested to enter the critical section then it performs the former, otherwise it takes the latter.

This algorithm satisfies ME1 due to the round robin approach taken meaning that two processes cannot execute in the critical section at the same time. There is only a single token that is passed round, therefore only one process executing in a critical section. ME2 is satisfied for similar reasons, as the token is being continuously passed around the ring each process will eventually gain access to the critical section (even in cases where it did not request it). ME3 is however not satisfied, as the clockwise rotation of the token goes against the ‘happened-before’ approach. A process p_N may make a request, but as the token is being passed through the ring another process p_{N-2} may similarly make a request. As the token reaches the second process first, even though the request was made afterwards, ME3 is not satisfied.

E. Multicast

The last algorithm discussed in this report, the multicast approach aims to address the issues of the previous two algorithms. Namely, it attempts to satisfy all three properties ME1, ME2 and ME3. Fundamentally, multicast works based on waiting for a response by all other processes to a request made by one process that wishes to enter the critical section. This is illustrated in Figure 3, with the connections between nodes on the graph representing the responses and requests made. To assist in this approach, Lamport clocks and timestamps are employed. It is assumed both that all processes have communication channels with one another and self-identify through the use of numerical IDs and timestamps. Each process has a *state*; either *WANTED*, *RELEASED* or *HELD*. These correspond to wanting the token, no longer needing the token and holding the token. When multiple processes issue a request close in time to each other, the one with the lowest Lamport timestamp is the first to collect replies and therefore first to execute in the critical section.

This algorithm satisfies ME1 as a process cannot execute in the critical section until receiving an approval by all other processes. Similarly, ME2 is satisfied as every process

eventually has the opportunity to execute in the critical section. With the assumption that interprocess communication is never broken, the usage of the *state* variable in conjunction with Lamport times assures that all processes have the opportunity to execute in the critical section. ME3 is satisfied as processes with the lowest Lamport timestamp receive the token before others, keeping the ‘happened-before’ approach. There is the possibility for situations arising where Lamport timestamps are equal, in which case fairness is enforced through the consistency that requests are ordered based on processes’ identifiers.




Fig. 1. An example of the *Central Server* algorithm




Fig. 2. The *Ring Token* algorithm

ACKNOWLEDGEMENT

This report acknowledges the teachings of Dr. Avinash Malik and Ms. Jesin James in the course COMPSYS 725: Distributed Cyber-Physical Systems taught at the University of Auckland in Semester Two of the year 2020.

REFERENCES

- [1] George F. Coulouris, Jean Dollimore, Tim Kindberg and Gordon Blair, “Distributed Systems: Concepts and Designs”, 5th ed, Boston, Massachusetts, Addison-Wesley; Pearson Education, 2011, pp. 41-49.

TABLE I
ALGORITHM COMPARISON

Algorithm	Properties			Metrics		
	Safety [ME1]	Liveness [ME2]	Fairness [ME3]	Consumed Bandwidth	Client Delay	Synchronisation Delay
Central Server	Satisfied	Satisfied	Not Satisfied	2 messages on <code>entry()</code> (first request, then token granted) with the process delayed by time taken for round-trip. EXIT takes a single release message	Client delay of 2; one <code>entry()</code> message and one <code>exit()</code> message with client only seeing a single message at a time	Time taken for a round-trip; release token to central server, then grant token to next process
Ring Token	Satisfied	Satisfied	Not Satisfied	Quite high; continuous consumption except for when a process executes inside the critical section	Client delay by process requesting critical section is between 0 and N messages (either it has received the token, or has just passed the token along). <code>exit()</code> requires only one message.	Anywhere from 1 to N messages.
Multicast	Satisfied	Satisfied	Satisfied	<code>entry()</code> takes $2(N - 1)$ messages; $N - 1$ for multicasting the request followed by $N - 1$ replies. Most expensive algorithm in terms of bandwidth.	Client delay is the round-trip time (<code>entry()</code> and <code>exit()</code>) ignoring delays incurred in multicasting the request message to all other processes.	The time taken for one message. Lowest synchronisation delay.




Fig. 3. An example of the *Multicast* algorithm